A

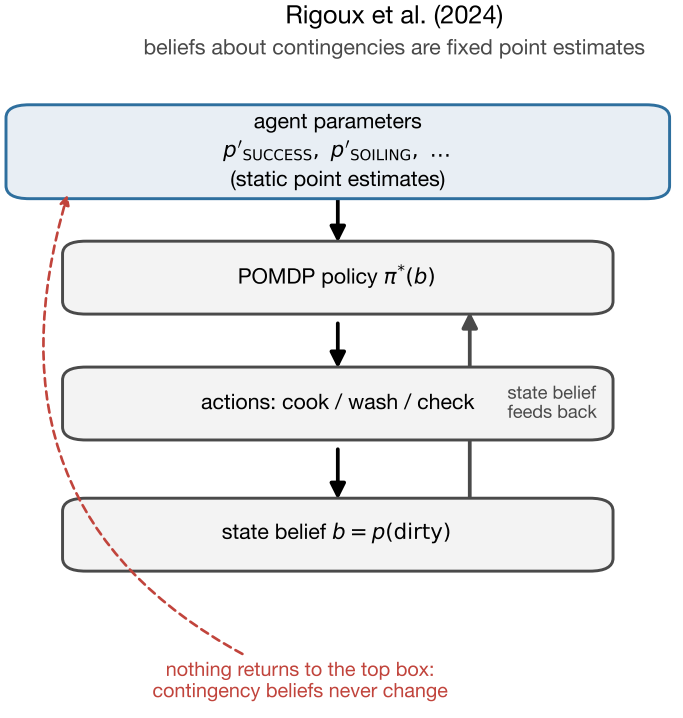
Rigoux et al. (2024)
beliefs about contingencies are fixed point estimates
agent parameters
p'_{SUCCESS}, p'_{SOILING}, …
(static point estimates)
POMDP policy $\pi^*(b)$
actions: cook / wash / check
state belief feeds back
state belief $b = p(\text{dirty})$
nothing returns to the top box:
contingency beliefs never change

B

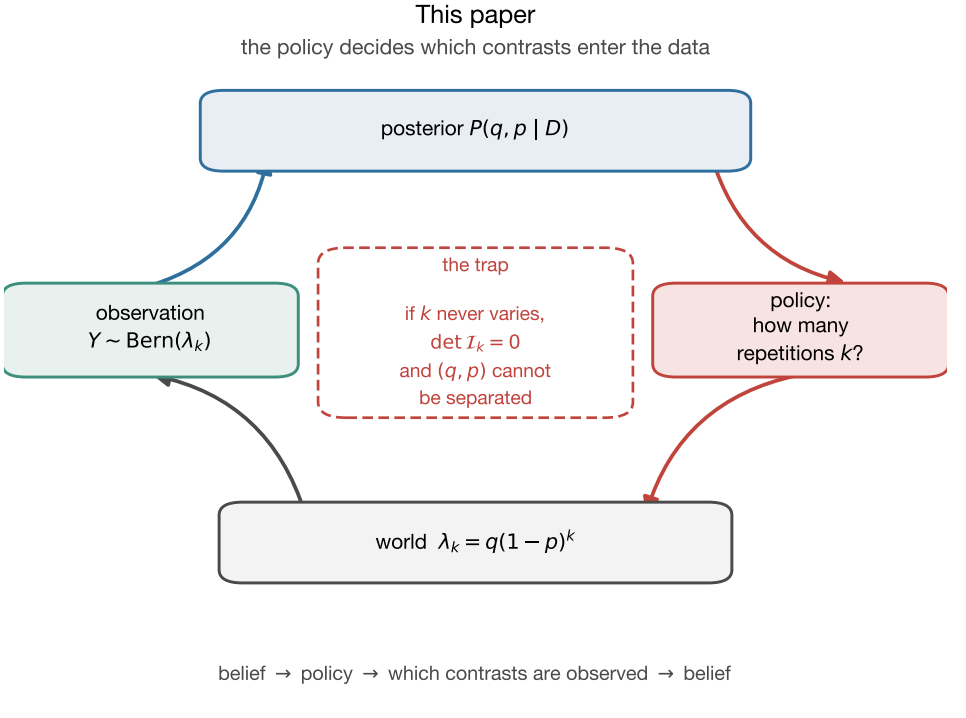
This paper
the policy decides which contrasts enter the data
posterior $P(q, p \mid D)$
observation
$Y \sim \text{Bern}(\lambda_k)$
the trap
if k never varies,
$\det \mathcal{I}_k = 0$
and (q, p) cannot
be separated
policy:
how many
repetitions k?
world $\lambda_k = q(1-p)^k$
belief → policy → which contrasts are observed → belief